

NNFL Project

November 19, 2018

1 Importing necessary libraries

```
In [95]: import skimage.data
import skimage.transform

import os
import numpy as np
import cv2
import pickle
from imutils import paths

from sklearn.preprocessing import LabelBinarizer
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report

from keras.models import Sequential
from keras.models import Model
from keras.optimizers import SGD
from keras.layers.normalization import BatchNormalization
from keras.layers.convolutional import Conv2D
from keras.layers.convolutional import MaxPooling2D
from keras.layers.core import Activation
from keras.layers.core import Flatten
from keras.layers.core import Dropout
from keras.layers.core import Dense
from keras.layers import Concatenate
from keras.preprocessing.image import ImageDataGenerator
from keras.engine.input_layer import Input
from keras.utils.vis_utils import plot_model
from keras import backend as K

# set the matplotlib backend so figures can be saved in the background
import matplotlib
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
matplotlib.use("Agg")
```

```
# Allow image embedding in notebook
%matplotlib inline
```

```
In [3]: #data set paths
BTS_ROOT = '/media/prithvi/DATA/acads/NNFL/assignment 2/BelgiumTS datasets/'
BTSC_TRAINING = os.path.join(BTS_ROOT,"BelgiumTSC_Training")
BTSC_TESTING= os.path.join(BTS_ROOT,"BelgiumTSC_Testing")
BTSC_TRAINING_SMALL = os.path.join(BTS_ROOT,"BelgiumTSC_Training_small")
```

2 Loading the Dataset

```
In [90]: #function to load data from data_dir
def load_data(data_dir):
    # Get all subdirectories of data_dir. Each represents a label.
    directories = [d for d in os.listdir(data_dir)
                    if os.path.isdir(os.path.join(data_dir, d))]

    # Loop through the label directories and collect the data in
    # two lists, labels and images.
    labels = []
    images = []
    for d in directories:
        label_dir = os.path.join(data_dir, d)
        file_names = [os.path.join(label_dir, f)
                       for f in os.listdir(label_dir)
                       if f.endswith(".ppm")]
        for f in file_names:
            images.append(skimage.data.imread(f))
            labels.append(int(d))
    return images, labels

In [92]: #loading the data
# images, labels = load_data(BTSC_TRAINING)
images, labels = load_data(BTSC_TRAINING)
```

2.1 Displaying an example image from each class

```
In [93]: def display_images_and_labels(images, labels):
    """Display the first image of each label."""
    unique_labels = set(labels)
    plt.figure(figsize=(15, 15))
    i = 1
    for label in unique_labels:
        # Pick the first image for each label.
        image = images[labels.index(label)]
        plt.subplot(8, 8, i) # A grid of 8 rows x 8 columns
        plt.axis('off')
        plt.title("Label {0} ({1})".format(label, labels.count(label)))
```

```

        i += 1
        _ = plt.imshow(image)
    plt.show()

```

```
display_images_and_labels(images, labels)
```



2.2 Resizing Images and One-Hot encoding class labels

```

In [7]: '''
        resizing, training data
        storing it as np array of floats normalised to interval [0,1]
        '''
        trainX = np.array([ cv2.resize(image,(48,48)) for image in images ] , dtype="float")/255

```

```

#converting labels to one hot encoding
lb= LabelBinarizer()
trainY = lb.fit_transform(labels) #finds unique labels and one hot encodes

#reading test data
testims,testLabels= load_data(BTSC_TESTING)
testX = np.array([ cv2.resize(image,(48,48)) for image in testims], dtype="float")/255
testY = lb.transform(testLabels) #transform() just encodes based on classes detected i

In [8]: print("No. of training examples = ", len(trainX))
        print(trainX[0].shape)
        print("No. of test examples = ",len(testX))
        print("No. of classes =",len(lb.classes_))

No. of training examples = 4575
(48, 48, 3)
No. of test examples = 2554
No. of classes = 62

```

3 Architecture

```

In [104]: model_from_paper = mpimg.imread("architecture.png")
          legend = skimage.data.imread("legend.png")
          plt.figure(figsize=(10,10),dpi=512)
          plt.subplot(1,2,1)
          plt.imshow(legend,aspect='equal')
          plt.axis("off");
          plt.subplot(1,2,2)
          plt.imshow(model_from_paper,aspect='auto')
          plt.axis('off');

```



```

#second convolutional layer
c2 = Conv2D(filters=128, kernel_size=(3,3), strides=1, padding="same")(p1)

#second pooling layer
p2 = MaxPooling2D(pool_size = (2,2), strides=2)(c2)
p2 = Dropout(0.25)(p2)

#third convolutional layer
c3 = Conv2D(filters=64, kernel_size=(3,3), strides=1, padding="same")(p2)

#third pooling layer
p3 = MaxPooling2D(pool_size = (2,2), strides=2)(c3)
p3 = Dropout(0.25)(p3)

#flattening outputs of p1,p2 and p3
f11 = Flatten()(p1)
f12 = Flatten()(p2)
f13 = Flatten()(p3)

#fully connected layers
fc1 = Dense(96)(f11)
fc1 = Dropout(0.25)(fc1)

fc2 = Dense(96)(f12)
fc2 = Dropout(0.25)(fc2)

fc3 = Dense(96)(f13)
fc3 = Dropout(0.25)(fc3)

#merging layer
m1 = Concatenate()([fc1, fc2, fc3])

#fully connected layer
fc4 = Dense(300, activation = 'relu')(m1)
fc4 = Dropout(0.25)(fc4)

#fully connected
fc5 = Dense(100,activation = 'relu')(fc4)
fc5 = Dropout(0.25)(fc5)

#fully connected (final layer - 62 outputs)
fc6 = Dense(len(lb.classes_), activation = 'softmax')(fc5)

#model
model = Model(inputs = ip, outputs = fc6)

# initialize our initial learning rate and # of epochs to train for

```

```
INIT_LR = 0.03
MOMENTUM = 0.3
```

```
#INIT_LR = 0.03 and MOMENTUM = 0.3 are the best settings found so far
```

```
# compile the model using SGD as our optimizer and categorical  
# cross-entropy loss (softmax loss)
```

```
opt = SGD(lr = INIT_LR, momentum = MOMENTUM)  
model.compile(loss="categorical_crossentropy", optimizer=opt, metrics=["accuracy"])
```

```
In [115]: print(model.summary())  
plot_model(model, show_shapes= True)
```

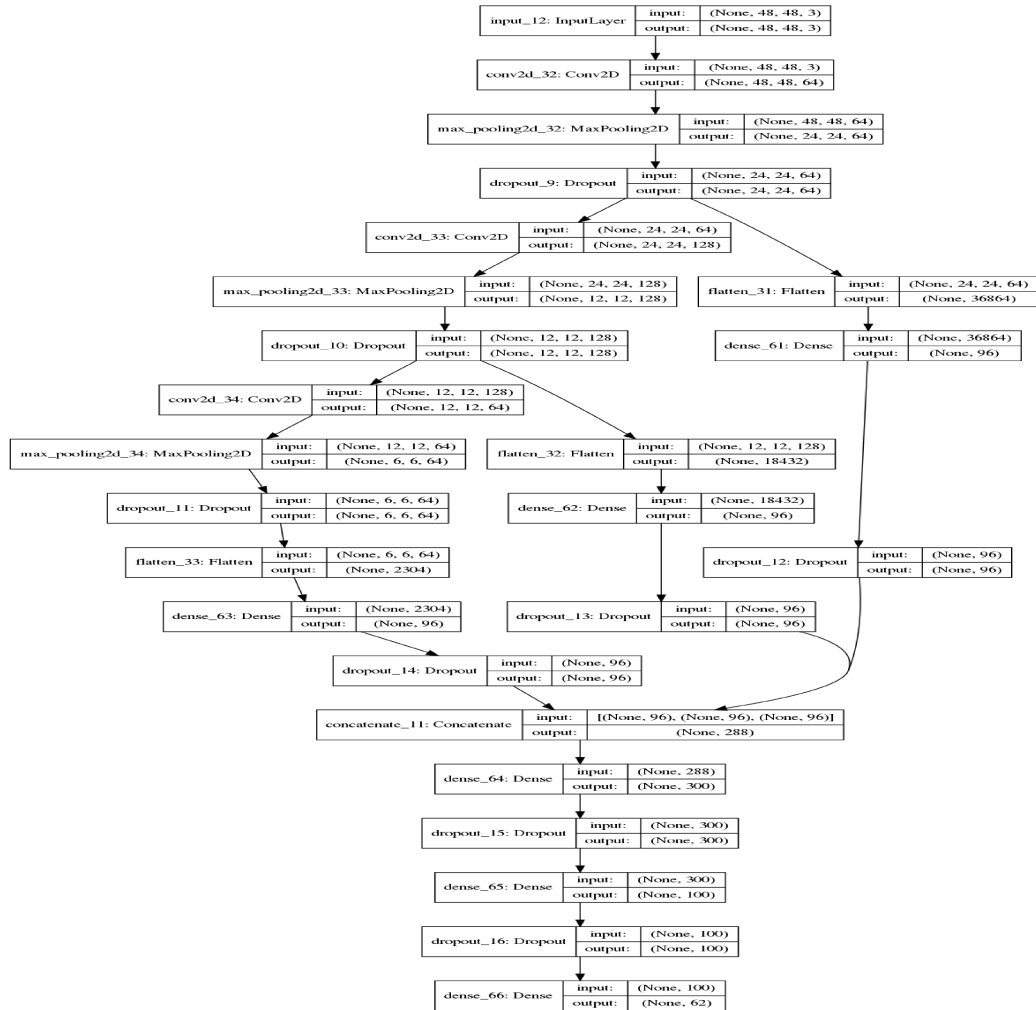
Layer (type)	Output Shape	Param #	Connected to
input_12 (InputLayer)	(None, 48, 48, 3)	0	
conv2d_32 (Conv2D)	(None, 48, 48, 64)	4864	input_12[0][0]
max_pooling2d_32 (MaxPooling2D)	(None, 24, 24, 64)	0	conv2d_32[0][0]
dropout_9 (Dropout)	(None, 24, 24, 64)	0	max_pooling2d_32[0][0]
conv2d_33 (Conv2D)	(None, 24, 24, 128)	73856	dropout_9[0][0]
max_pooling2d_33 (MaxPooling2D)	(None, 12, 12, 128)	0	conv2d_33[0][0]
dropout_10 (Dropout)	(None, 12, 12, 128)	0	max_pooling2d_33[0][0]
conv2d_34 (Conv2D)	(None, 12, 12, 64)	73792	dropout_10[0][0]
max_pooling2d_34 (MaxPooling2D)	(None, 6, 6, 64)	0	conv2d_34[0][0]
dropout_11 (Dropout)	(None, 6, 6, 64)	0	max_pooling2d_34[0][0]
flatten_31 (Flatten)	(None, 36864)	0	dropout_9[0][0]
flatten_32 (Flatten)	(None, 18432)	0	dropout_10[0][0]
flatten_33 (Flatten)	(None, 2304)	0	dropout_11[0][0]
dense_61 (Dense)	(None, 96)	3539040	flatten_31[0][0]
dense_62 (Dense)	(None, 96)	1769568	flatten_32[0][0]
dense_63 (Dense)	(None, 96)	221280	flatten_33[0][0]

dropout_12 (Dropout)	(None, 96)	0	dense_61[0][0]
dropout_13 (Dropout)	(None, 96)	0	dense_62[0][0]
dropout_14 (Dropout)	(None, 96)	0	dense_63[0][0]
concatenate_11 (Concatenate)	(None, 288)	0	dropout_12[0][0] dropout_13[0][0] dropout_14[0][0]
dense_64 (Dense)	(None, 300)	86700	concatenate_11[0][0]
dropout_15 (Dropout)	(None, 300)	0	dense_64[0][0]
dense_65 (Dense)	(None, 100)	30100	dropout_15[0][0]
dropout_16 (Dropout)	(None, 100)	0	dense_65[0][0]
dense_66 (Dense)	(None, 62)	6262	dropout_16[0][0]
=====			
Total params: 5,805,462			
Trainable params: 5,805,462			
Non-trainable params: 0			

None			

4.2 Architecture of the network made

```
In [116]: model_diagram=skimage.data.imread("model.png")
plt.figure(figsize=(10,10),dpi=512)
plt.imshow(model_diagram,aspect='auto')
plt.axis("off");
```

5 Training

In [119]: *#train the model*

EPOCHS = 50

BATCH_SIZE = 32

```
gen = ImageDataGenerator(width_shift_range = 0.1, height_shift_range = 0.1, zoom_range = 0.1)
H = model.fit_generator(gen.flow(trainX,trainY), validation_data=(testX,testY), steps_per_epoch=1000)
```

Epoch 1/50

142/142 [=====] - 49s 345ms/step - loss: 0.3526 - acc: 0.8875 - val_loss: 0.3829 - val_acc: 0.8818

Epoch 2/50

142/142 [=====] - 50s 350ms/step - loss: 0.3829 - acc: 0.8818 - val_loss: 0.3829 - val_acc: 0.8818

Epoch 3/50

```

142/142 [=====] - 49s 346ms/step - loss: 0.3714 - acc: 0.8840 - val_1
Epoch 4/50
142/142 [=====] - 49s 346ms/step - loss: 0.3412 - acc: 0.8886 - val_1
Epoch 5/50
142/142 [=====] - 52s 364ms/step - loss: 0.3184 - acc: 0.8965 - val_1
Epoch 6/50
142/142 [=====] - 52s 367ms/step - loss: 0.3229 - acc: 0.8961 - val_1
Epoch 7/50
142/142 [=====] - 49s 344ms/step - loss: 0.3218 - acc: 0.8998 - val_1
Epoch 8/50
142/142 [=====] - 52s 364ms/step - loss: 0.3202 - acc: 0.8994 - val_1
Epoch 9/50
142/142 [=====] - 48s 342ms/step - loss: 0.3202 - acc: 0.8998 - val_1
Epoch 10/50
142/142 [=====] - 49s 342ms/step - loss: 0.2935 - acc: 0.9087 - val_1
Epoch 11/50
142/142 [=====] - 49s 347ms/step - loss: 0.2887 - acc: 0.9109 - val_1
Epoch 12/50
142/142 [=====] - 50s 353ms/step - loss: 0.2899 - acc: 0.9062 - val_1
Epoch 13/50
142/142 [=====] - 51s 358ms/step - loss: 0.2582 - acc: 0.9170 - val_1
Epoch 14/50
142/142 [=====] - 49s 345ms/step - loss: 0.2752 - acc: 0.9135 - val_1
Epoch 15/50
142/142 [=====] - 49s 346ms/step - loss: 0.2796 - acc: 0.9155 - val_1
Epoch 16/50
142/142 [=====] - 49s 346ms/step - loss: 0.2620 - acc: 0.9183 - val_1
Epoch 17/50
142/142 [=====] - 49s 346ms/step - loss: 0.2381 - acc: 0.9309 - val_1
Epoch 18/50
142/142 [=====] - 49s 347ms/step - loss: 0.2468 - acc: 0.9223 - val_1
Epoch 19/50
142/142 [=====] - 49s 347ms/step - loss: 0.2344 - acc: 0.9258 - val_1
Epoch 20/50
142/142 [=====] - 54s 382ms/step - loss: 0.2228 - acc: 0.9289 - val_1
Epoch 21/50
142/142 [=====] - 50s 350ms/step - loss: 0.2194 - acc: 0.9238 - val_1
Epoch 22/50
142/142 [=====] - 49s 347ms/step - loss: 0.2386 - acc: 0.9252 - val_1
Epoch 23/50
142/142 [=====] - 50s 349ms/step - loss: 0.2170 - acc: 0.9344 - val_1
Epoch 24/50
142/142 [=====] - 51s 358ms/step - loss: 0.2308 - acc: 0.9276 - val_1
Epoch 25/50
142/142 [=====] - 50s 349ms/step - loss: 0.2321 - acc: 0.9269 - val_1
Epoch 26/50
142/142 [=====] - 49s 346ms/step - loss: 0.2009 - acc: 0.9351 - val_1
Epoch 27/50

```

```

142/142 [=====] - 49s 346ms/step - loss: 0.2120 - acc: 0.9344 - val_loss: 0.2120
Epoch 28/50
142/142 [=====] - 49s 348ms/step - loss: 0.2283 - acc: 0.9300 - val_loss: 0.2283
Epoch 29/50
142/142 [=====] - 49s 347ms/step - loss: 0.2115 - acc: 0.9333 - val_loss: 0.2115
Epoch 30/50
142/142 [=====] - 50s 353ms/step - loss: 0.1902 - acc: 0.9428 - val_loss: 0.1902
Epoch 31/50
142/142 [=====] - 50s 352ms/step - loss: 0.1892 - acc: 0.9414 - val_loss: 0.1892
Epoch 32/50
142/142 [=====] - 50s 349ms/step - loss: 0.1887 - acc: 0.9428 - val_loss: 0.1887
Epoch 33/50
142/142 [=====] - 50s 350ms/step - loss: 0.1837 - acc: 0.9410 - val_loss: 0.1837
Epoch 34/50
142/142 [=====] - 49s 344ms/step - loss: 0.1841 - acc: 0.9397 - val_loss: 0.1841
Epoch 35/50
142/142 [=====] - 50s 349ms/step - loss: 0.2043 - acc: 0.9423 - val_loss: 0.2043
Epoch 36/50
142/142 [=====] - 51s 356ms/step - loss: 0.1529 - acc: 0.9555 - val_loss: 0.1529
Epoch 37/50
142/142 [=====] - 49s 347ms/step - loss: 0.1718 - acc: 0.9472 - val_loss: 0.1718
Epoch 38/50
142/142 [=====] - 49s 348ms/step - loss: 0.1933 - acc: 0.9401 - val_loss: 0.1933
Epoch 39/50
142/142 [=====] - 50s 354ms/step - loss: 0.1804 - acc: 0.9467 - val_loss: 0.1804
Epoch 40/50
142/142 [=====] - 53s 370ms/step - loss: 0.1772 - acc: 0.9434 - val_loss: 0.1772
Epoch 41/50
142/142 [=====] - 53s 375ms/step - loss: 0.1841 - acc: 0.9417 - val_loss: 0.1841
Epoch 42/50
142/142 [=====] - 52s 367ms/step - loss: 0.1771 - acc: 0.9448 - val_loss: 0.1771
Epoch 43/50
142/142 [=====] - 52s 367ms/step - loss: 0.1640 - acc: 0.9478 - val_loss: 0.1640
Epoch 44/50
142/142 [=====] - 53s 371ms/step - loss: 0.1641 - acc: 0.9503 - val_loss: 0.1641
Epoch 45/50
142/142 [=====] - 53s 377ms/step - loss: 0.1715 - acc: 0.9465 - val_loss: 0.1715
Epoch 46/50
142/142 [=====] - 54s 379ms/step - loss: 0.1729 - acc: 0.9474 - val_loss: 0.1729
Epoch 47/50
142/142 [=====] - 52s 369ms/step - loss: 0.1724 - acc: 0.9459 - val_loss: 0.1724
Epoch 48/50
142/142 [=====] - 53s 370ms/step - loss: 0.1611 - acc: 0.9520 - val_loss: 0.1611
Epoch 49/50
142/142 [=====] - 51s 361ms/step - loss: 0.1539 - acc: 0.9527 - val_loss: 0.1539
Epoch 50/50
142/142 [=====] - 49s 345ms/step - loss: 0.1489 - acc: 0.9529 - val_loss: 0.1489

```

```
In [122]: # save the model and label binarizer to disk
model.save("model_dropout25")
f = open("binarizer", "wb")
f.write(pickle.dumps(lb))
f.close()
```

6 Testing

```
In [121]: # evaluate the network
predictions = model.predict(testX, batch_size=32)
print(classification_report(testY.argmax(axis=1), predictions.argmax(axis=1), labels=
```

	precision	recall	f1-score	support
0	0.86	1.00	0.92	6
1	0.96	1.00	0.98	27
2	1.00	0.71	0.83	7
3	0.75	0.50	0.60	6
4	0.92	0.92	0.92	12
5	1.00	0.67	0.80	3
6	0.62	0.83	0.71	6
7	0.99	1.00	0.99	90
8	0.63	1.00	0.77	12
9	1.00	1.00	1.00	7
10	1.00	0.96	0.98	28
11	1.00	1.00	1.00	4
12	1.00	0.67	0.80	3
13	1.00	0.90	0.95	39
14	0.62	1.00	0.77	15
15	1.00	1.00	1.00	5
16	1.00	0.25	0.40	12
17	0.98	0.99	0.99	183
18	1.00	0.98	0.99	122
19	1.00	1.00	1.00	163
20	1.00	1.00	1.00	3
21	0.98	0.98	0.98	45
22	0.98	1.00	0.99	61
23	1.00	0.80	0.89	15
24	0.91	0.77	0.83	13
25	0.60	1.00	0.75	3
26	0.50	1.00	0.67	2
27	1.00	1.00	1.00	9
28	0.91	1.00	0.95	51
29	0.93	1.00	0.97	28
30	0.97	1.00	0.99	37
31	1.00	1.00	1.00	86
32	1.00	0.99	1.00	422

33	1.00	1.00	1.00	3
34	1.00	1.00	1.00	9
35	1.00	0.95	0.97	154
36	1.00	1.00	1.00	4
37	1.00	1.00	1.00	31
38	0.97	1.00	0.99	213
39	1.00	0.99	0.99	99
40	1.00	0.96	0.98	48
41	1.00	1.00	1.00	11
42	0.64	1.00	0.78	9
43	1.00	0.17	0.29	6
44	1.00	1.00	1.00	3
45	0.96	0.92	0.94	84
46	0.43	0.50	0.46	6
47	0.86	1.00	0.93	31
48	1.00	1.00	1.00	3
49	1.00	0.33	0.50	3
50	1.00	1.00	1.00	3
51	1.00	1.00	1.00	3
52	0.75	1.00	0.86	3
53	1.00	1.00	1.00	24
54	1.00	1.00	1.00	48
55	1.00	0.93	0.97	15
56	0.97	1.00	0.99	33
57	1.00	0.78	0.88	41
58	0.90	1.00	0.95	9
59	0.68	1.00	0.81	17
60	1.00	0.55	0.71	11
61	0.95	1.00	0.98	105
micro avg	0.97	0.97	0.97	2554
macro avg	0.92	0.90	0.89	2554
weighted avg	0.97	0.97	0.97	2554

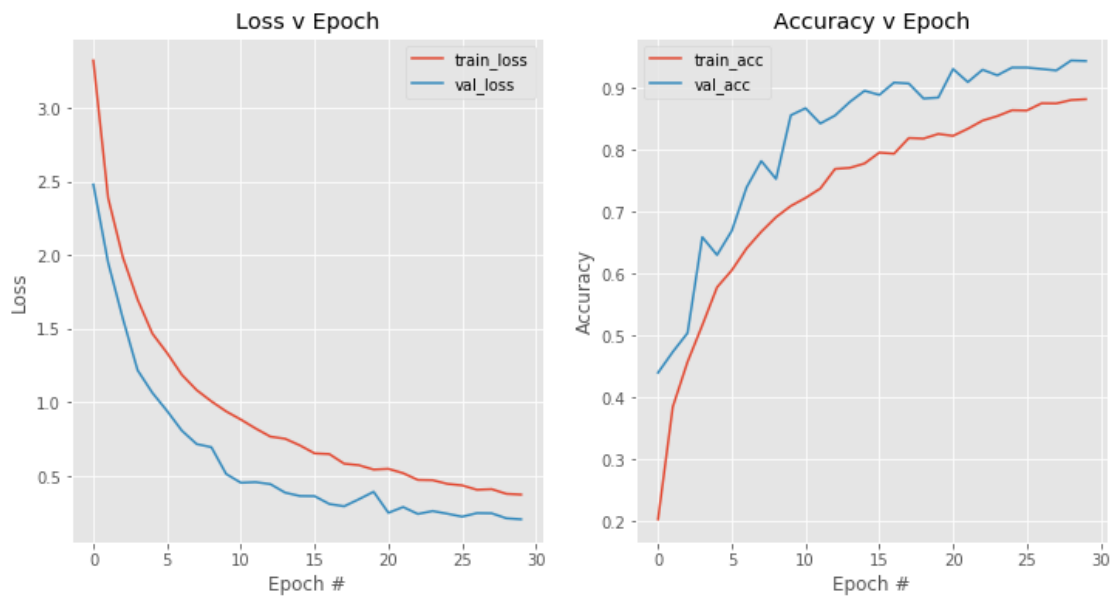
6.1 Plotting traing curves

```
In [118]: # plot the training loss and accuracy
N = np.arange(0, EPOCHS)
plt.style.use("ggplot")
plt.figure(figsize=(12,6))
plt.title("Training Loss and Accuracy (Simple NN)")
plt.subplot(1,2,1)
plt.title("Loss v Epoch")
plt.plot(N, H.history["loss"], label="train_loss")
plt.plot(N, H.history["val_loss"], label="val_loss")
plt.xlabel("Epoch #")
```

```

plt.ylabel("Loss")
plt.legend()
plt.subplot(1,2,2)
plt.title("Accuracy v Epoch")
plt.plot(N, H.history["acc"], label="train_acc")
plt.plot(N, H.history["val_acc"], label="val_acc")
plt.xlabel("Epoch #")
plt.ylabel("Accuracy")
plt.legend()
plt.savefig(".")

```



In []: